

Parallel Bat Optimization on GPU using CUDA

Jean Carlo Machado
Univerisdade do Estado de Santa Catarina
Mestrado de Computao Aplicada
UDESC
Santa Catarina, Joinville,
Email: contato@jeancarlomachado.com.br

Rafael Stubs Parpinelli
Univerisdade do Estado de Santa Catarina
UDESC
Santa Catarina, Joinville

Abstract—The ever increasing parallel processing power of GPU's is an compelling reason to implement performance demanding algorithms, like optimization techniques, using this technology. This work aimed to develop the bat algorithm on GPU and benchmark it against a CPU version. A set of experiments where conducted in order to measure the speedup difference. The results suggests that the GPU version is able to achieve relevant speedups in highly populational problems but for simpler cases the CPU version might outperform the GPU.

I. INTRODUCTION

With the aid of parallel computing, specially GPU computing, it's possible to tackle ever increasing computational problems in decreasing amounts of time. And CPU's are lagging behind GPU's in many aspects. To give a point of contact, [11] mentions that the ratio between many-core GPU's and multi-core CPU's for peak floating-point calculation throughput is about 10 to 1.

Meta-heuristics are a traditional method for solving complex problems. A promising meta-heuristics group for parallel computing is the swarm intelligence. Which encompasses many bio-inspirations like PSO, ACO, and the relatively novel BAT algorithm. Parpinelli [9], when referring to swarm intelligence algorithms, define that they are:

compounded by a distributed society/population of individuals where the control is also distributed among the individuals (there is no centralized control); and the individuals decision-making is stochastic and based only on local information.

The fact that swarm intelligence algorithms have composed by independent parts enables them to run in parallel. This work attempts to investigate the applicability of a relatively new and promising swarm algorithm: the BAT algorithm using the CUDA model.

Previously some demonstrations of the bat algorithm parallelized on CPU were presented in [7] and [6]. A recent publication about the use of the bat optimization can be found on [10]. However, til the moment of this publications there's no application of the bat on GPU tested against standard functions.

The rest of the paper is organized as following: Section II gives a overview of CUDA and its components. Section III studies the bat meta-heuristic, it's design and details. Section

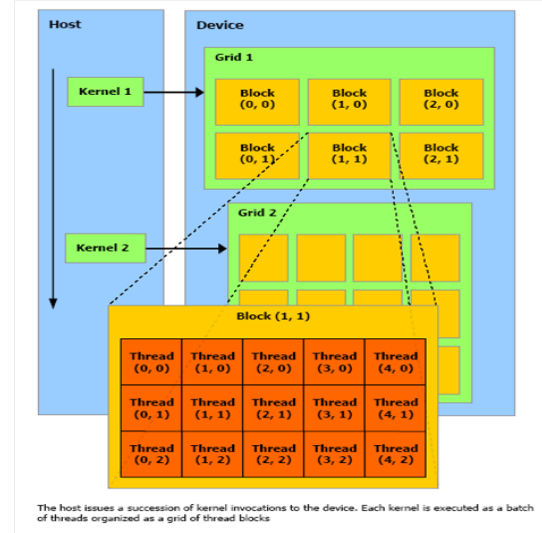


Fig. 1. The CUDA structural organization

IV details the experiments performed. Section V analyses the results.

II. CUDA

Compute Unified Device Architecture (CUDA) is a general purpose parallel computing programming model for parallel computations [3]. CUDA uses a SIMD *single data multiple execution* approach [8], where the concurrent code executes the same instruction but with divergent data.

The basic element of work in a CUDA device is a thread. That is aggruped and runs in parallel with 32 other threads called the warp. Wraps are grouped in blocks.

Streaming multiprocessors μ organized in blocks of streaming processors

A CUDA program starts by allocating resources on the GPU and dispatching work to de done on the GPU.

Once the GPU finishes it's work the GPU they are moved back from GPU memory to the CPU one. And the results can be returned.

For the developer, the basic element of concern is the kernel, which is the function executed on the GPU. A kernel

function defines the code to be executed by each of the massive numbers of threads to be invoked [8].

A. Performance concerns in CUDA

Since the CUDA model is complex, it's not trivial to setup an ideal configuration of a problem in the GPU. However many good practises exists and great speedup is already reported. Previous researches suggests that highly parallel applications may speedup up to 450 times [8].

Below are described some of the major concern while developing parallel algorithms in CUDA.

- When attempting to achieve an application's maximum performance, the primary concern often is managing global memory latency [8]. It's preferable to use local memory and shared memory instead of global since they are much faster.
- In an ideal enviroment when the GPU starts working the less the messaging passing between the GPU and the CPU, the better.
- Decrease data interchange between the host and the GPU to the mininum.
- Prefer floats to doubles when possible. The older versions of CUDA GPU's had serious bottlenets using large variables type, but this is a issues disappearing

III. BAT ALGORITHM

The bat algorithm is a populational meta-heuristic introduced by Yang [1] in 2010. It uses the inspiration of micro-bats which uses a type of sonar, called echolocation, to detect prey, avoid obstacles, and locate their roosting crevices in the dark [1].

The bat algorithm has two parameters: the pulse rate and the loudness. As time goes by the pulse-rate tends to increase and the loudness to decrease. The inspiration comes from the behavior of some bats that use slow and loud pitches while in search for a prey and quick and low pitches when in persecution of one. In the bat algorithm the loudness is used as a way of accepting bad results (diversification) and pulse rate as a way of selecting local search (exploitation).

As the base algorithm we used the bat as proposed by [2]. Since this paper has a more detailed implementation than the original one. The algorithm can be found in Figure 2.

Some distinctions of the original paper are worth noticing:

The selection of new results on the original paper tends to be more greedy. On the original paper, for accepting new results on each iteration, the loudness must be greater than a aleatory number and the fitness of the candidate must be better than the current best. However on the version proposed by [2] its used an *OR* operator, so more candidates are accepted.

The or operator tends to explore the search space better (more diversity). Another divergent aspect of the alternative bat is that it contains a distortion of a single dimension of the results, in order to further increase the diversity.

At last, the random algorithm used was not specified on any of the papers, nevertheless on our simulations the Xorshift algorithm was used.

```

1: Parameters :  $n, \alpha, \lambda$ 
2: initialize bats
3: evaluate fitness
4: selects best  $\vec{x}_*$ 
5: while stop criteria false do
6:   for each bat do
7:      $f_i = f_{min} + (f_{max} - f_{min})\beta, \beta \in \beta[0, 1]$ 
8:      $\vec{v}_i^{t+1} = \vec{v}_i^t + (\vec{x}_i^t + \vec{x}_*^t)f_i$ 
9:      $\vec{x}_{temp} = \vec{x}_i^t + \vec{v}_i^{t+1}$ 
10:    if  $rand < r_i, rand \in [0, 1]$  then
11:       $\vec{x}_{temp} = \vec{x}_* + \epsilon A_m, \epsilon \in [-1, 1]$ 
12:    end if
13:    single dimension perturbation in  $x_{temp}$ 
14:    if  $a < A_i$  or  $f(\vec{x}_{temp}) \leq f(\vec{x}_i), a \in [0, 1]$  then
15:       $\vec{x}_i^t = \vec{x}_{temp}$ 
16:       $r_i = exp(\lambda * i)$ 
17:       $A_i = A_0 * \alpha^i$ 
18:    end if
19:  end for
20:  selects best  $\vec{x}_*$ 
21: end while

```

Fig. 2. Pseudo-code CPU




Fig. 3. GPU process flow

A. Bat Design on GPU

A convenient approach to model the swarm in the GPU is to use each thread as an individual. [4] used a similar method for a GPU implementation for the PSO algorithm, and he says

Each element of the particle is treated individually in the thread, allowing efficient data parallelism whitout risk of starvation or race conditions.

In the bat algorithm synchronization must occur on the selection of the best individual of the iteration. The best individual is kept in the threaded memory of the GPU which

```

1: Parameters :  $n, \alpha, \lambda$ 
2: initialize bats asynchronously
3: evaluate fitness
4: synchronize threads
5: selects best  $\vec{x}_*$ 
6: while stop criteria false do
7:   for each thread do
8:      $f_i = f_{min} + (f_{max} - f_{min})\beta, \beta \in \beta[0, 1]$ 
9:      $\vec{v}_i^{t+1} = \vec{v}_i^t + (\vec{x}_i^t + \vec{x}_*^t)f_i$ 
10:     $\vec{x}_{temp} = \vec{x}_i^t + \vec{v}_i^{t+1}$ 
11:    if  $rand < r_i, rand \in [0, 1]$  then
12:       $\vec{x}_{temp} = \vec{x}_* + \epsilon A_m, \epsilon \in [-1, 1]$ 
13:    end if
14:    single dimension perturbation in  $x_{temp}$ 
15:    if  $a < A_i^t$  or  $f(\vec{x}_{temp}) \leq f(\vec{x}_i), a \in [0, 1]$  then
16:       $\vec{x}_i^t = \vec{x}_{temp}$ 
17:       $r_i = exp(\lambda * i)$ 
18:       $A_i = A_0 * \alpha^i$ 
19:    end if
20:  end for
21:  synchronize threads
22:  selects best  $\vec{x}_*$ 
23: end while

```

Fig. 4. Pseudo-code GPU

has a limit of 16KB, probably not feasible for more complex problems.

The initial random number generator used on the GPU was the MTGP32. However later was discovered it's not recommended to use more than 256 threads per block with it [5]. So we moved the algorithms to the CUDA default XORWOR.

IV. EXPERIMENTS

For accessing the performance of the algorithm a set of experiments where elaborated as follows.

The benchmark functions used were the following:

- Ackley
- Griewank
- Rastrigin
- Rosenbrock

The experiments were executed on a machine with the following configuration:

Intel(R) Core(TM) i5-4460 CPU @ 3.20GHz
GK208 GeForce GT 720 1024 MB of vram
Compute capability 3.5
Kepler GM10x

Each experiment was executed a total of 20 times. A total of 10 thousand iterations where performed in each experiment. The benchmark functions were all normalized to work with 100 dimensions for each test. Table I details each experiment performed.

TABLE I
EXPERIMENTS

Name	Function	Dimensions	Agents
E1	Ackley	100	256
E2	Ackley	100	768
E3	Griewank	100	256
E4	Griewank	100	768
E5	Rastrigin	100	256
E6	Rastrigin	100	768
E7	Rosenbrock	100	256
E8	Rosenbrock	100	768




Fig. 5. Convergence on Rosenbrock

V. RESULTS

In this section are described the speedup and convergence results.

The fitness of almost all functions presented a slight worse result when compared with the CPU version. Since there's no noticeable difference in the design we assume that's related to the way the GPU process random numbers.

Anyway the purpose of this research is to focus on the speedups so it seems like a reasonable trade off.

The results shows that, as the bat population increases the performance increases as well on the GPU. The contrary is observed for the CPU.

TABLE II
CPU RESULTS

Time Avg	Time SD	Fit Avg	Fit SD
(E1) 49.4428	0.0557314	4.44089e-16	2.05196e-22
(E2) 161.439	0.155131	4.44089e-16	2.82843e-22
(E3) 61.3661	5.06578	0	0
(E4) 162.761	57.8119	0	0
(E5) 52.0624	9.25666	0	0
(E6) 171.986	14.9089	0	0
(E7) 20.4486	0.0847218	98.9875	0.0405622
(E8) 74.3533	0.186482	98.9864	0.0204378

TABLE III
GPU RESULTS

Time Avg	Time SD	Fit Avg	Fit SD
(E1) 17.2255	0.708198	12.8881	2.40027
(E2) 10.9591	0.23902	10.9412	3.23942
(E3) 24.2459	0.740923	2.04281e-15	2.60744e-16
(E4) 15.0012	2.01505e-15	2.55402e-16	0.0394986
(E5) 30.4483	2.0005	0	0
(E6) 14.4247	0.0543432	0	0
(E7) 28.9867	1.24554	105.03	31.2888
(E8) 15.4403	0.326284	101.793	140.393

VI. CONCLUSION

With this work it's clear that is possible to speedup the bat metaheuristic using GPU. Notwithstanding the best results are only achievable on really complex problems with many dimensions.

VII. FURTHER WORKS

Since the CPU version developed was single threaded it may had some disadvantages in speedup. The advantages of the algorithm may be tested against a threaded CPU implementation.

As the amount of memory is limited in the threaded model. In the future it may be explored the usage of blocks as representation for the dimensions in which each bat details.

A sub-population approach may also work, considering each GPU block as it's boundaries, somewhat similar to the work made on parallel bat on CPU by [6].

REFERENCES

- [1] X. S. Yang, "A New Metaheuristic Bat-Inspired Algorithm", in Studies in Computational Intelligence, Springer Berlin, 284, 2010
- [2] J.A. Cordeiro and R.S. Parpinelli and H.S. Lopes, "Análise de Sensibilidade dos Parâmetros do Bat Algorithm e Comparação de Desempenho," Department of Bioinformatics, UTFPR
- [3] NVIDIA, "CUDA C Programming Guide," no. July. NVIDIA Corporation, 2013
- [4] D. L. Souza et al., "PSO-GPU: Accelerating Particle Swarm Optimization in CUDA-Based Graphics Processing Units," in Laboratório de Computação Natural (CESUPA)
- [5] NVIDIA. Bit Generation with the MTGP32 generator [Online]. Available: <http://docs.nvidia.com/cuda/curand/device-api-overview.html>
- [6] C. Tsai, et al. "Parallelized Bat Algorithm with a Communication Strategy," in Modern Advances in Applied Intelligence, 27th International Conference of Industrial Engineering and Other Applications of Applied Intelligent Systems, IEA/AIE 2014
- [7] T. Dao et al., "Parallel bat algorithm for optimizing makespan in job scheduling problems", in Springer Science Review, 2015
- [8] S. Ryoo et al., "Optimization Principles and Application Performance Evaluation on a Multithreaded GPU using CUDA," in Center of Reliable and High-Performance Computing, University of Illinois
- [9] Parpinelli, R.S. and Lopes, H.S. (2011) 'New inspirations in swarm intelligence: a survey', Int. J. Bio-Inspired Computation, Vol. 3, No. 1 pp.1-16
- [10] A. R. Choudhury, "A GPU Implementation of a Bat Algorithm Trained Neural Network," in Neural Information Processing 23rd International Conference, ICONIP 2016
- [11] W. W. Hwu and D. Kirk, "Programming massively parallel processors,," in Springer Verlag GmbH, 2010